



15CSE312

COMPUTER NETWORKS

3-0-0 3

Amrita Vishwa Vidyapeetham
Amritapuri Campus





Chapter 5: Data Link Layer

LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

All material copyright 1996-2016

J.F Kurose and K.W. Ross, All Rights Reserved

Data Link Layer- Introduction

This layer is the protocol layer that **transfers data between nodes** on a network segment across the physical layer. The data link layer provides the functional and procedural means to transfer data between network entities and may also provide the means to detect and possibly correct errors that can occur in the physical layer.

Any device that runs a link-layer protocol as a node. **Nodes** include hosts, routers, switches, and WiFi access points We will also refer to the communication channels that connect adjacent nodes along the communication path as **links**. In order for a datagram to be transferred from source host to destination host, it must be moved over each of the individual links in the end-to-end path.

The basic service of any link layer is to move a datagram from one node to an adjacent node over a single communication link

Link layer, LANs: outline

5.1 introduction, services

5.2 error detection, correction

5.3 multiple access protocols

5.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

Switched Local Area Networks

- 2 servers, 1 router with 4 switches
- Because these switches operate at the link layer, they switch link-layer frames (rather than network-layer datagrams), don't recognize network-layer addresses
- Instead of using IP addresses use link-layer addresses to forward link layer frames through the network of switches
- Hosts and routers have link-layer addresses
- **Address Resolution Protocol (ARP)**, which provides a mechanism to translate IP addresses to link-layer addresses.

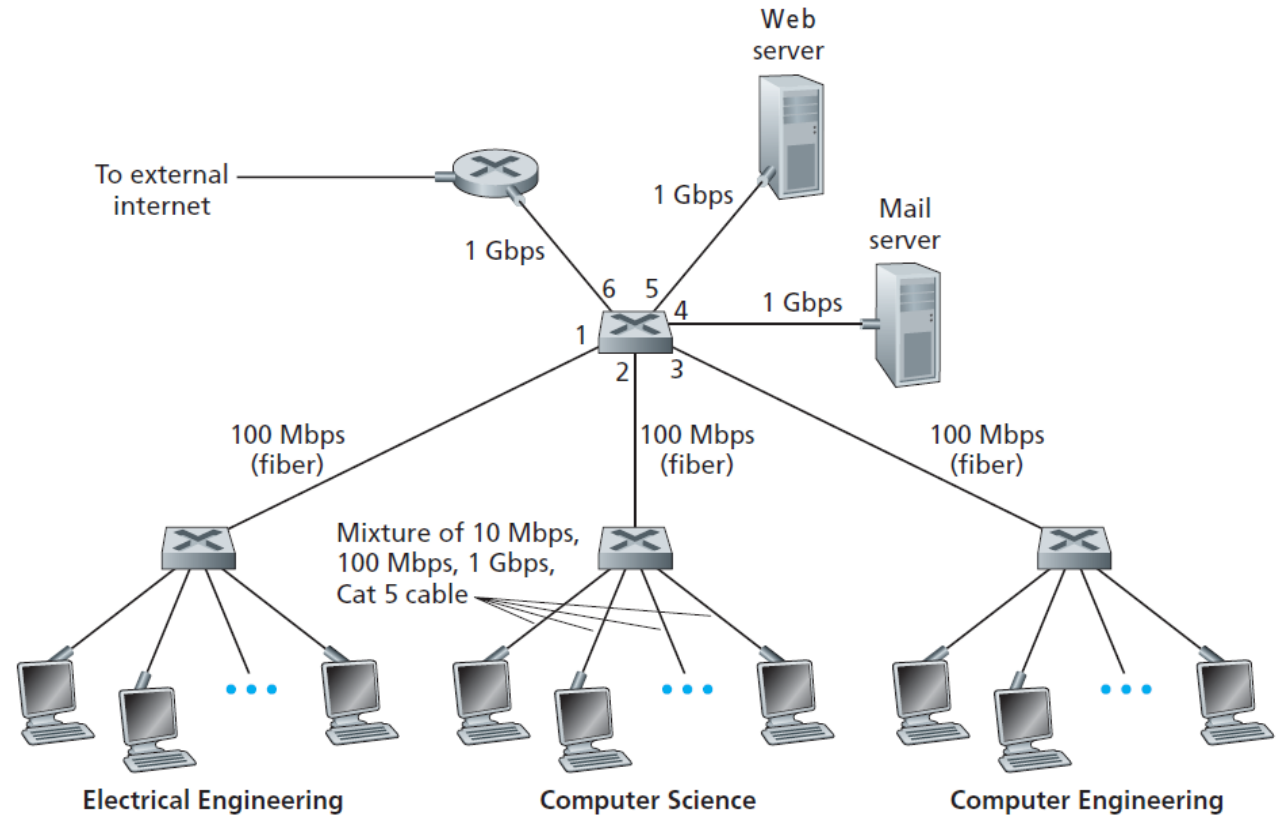


Figure 5.15 ♦ An institutional network connected together by four switches

MAC addresses and ARP (Address Resolution protocol)

- 32/128-bit IP address:
 - *network-layer* address for interface
 - used for **layer 3** (network layer) forwarding
- **MAC** (or LAN or physical or Ethernet) address:
 - function: *used ‘locally’ to get frame from one interface to another physically-connected interface (same network, in IP-addressing sense)*
 - **48 bit MAC address** (for most LANs) burned in NIC ROM, also sometimes software settable
 - e.g.: 1A-2F-BB-76-09-AD ——— hexadecimal (base 16) notation
(each “number” represents 4 bits)

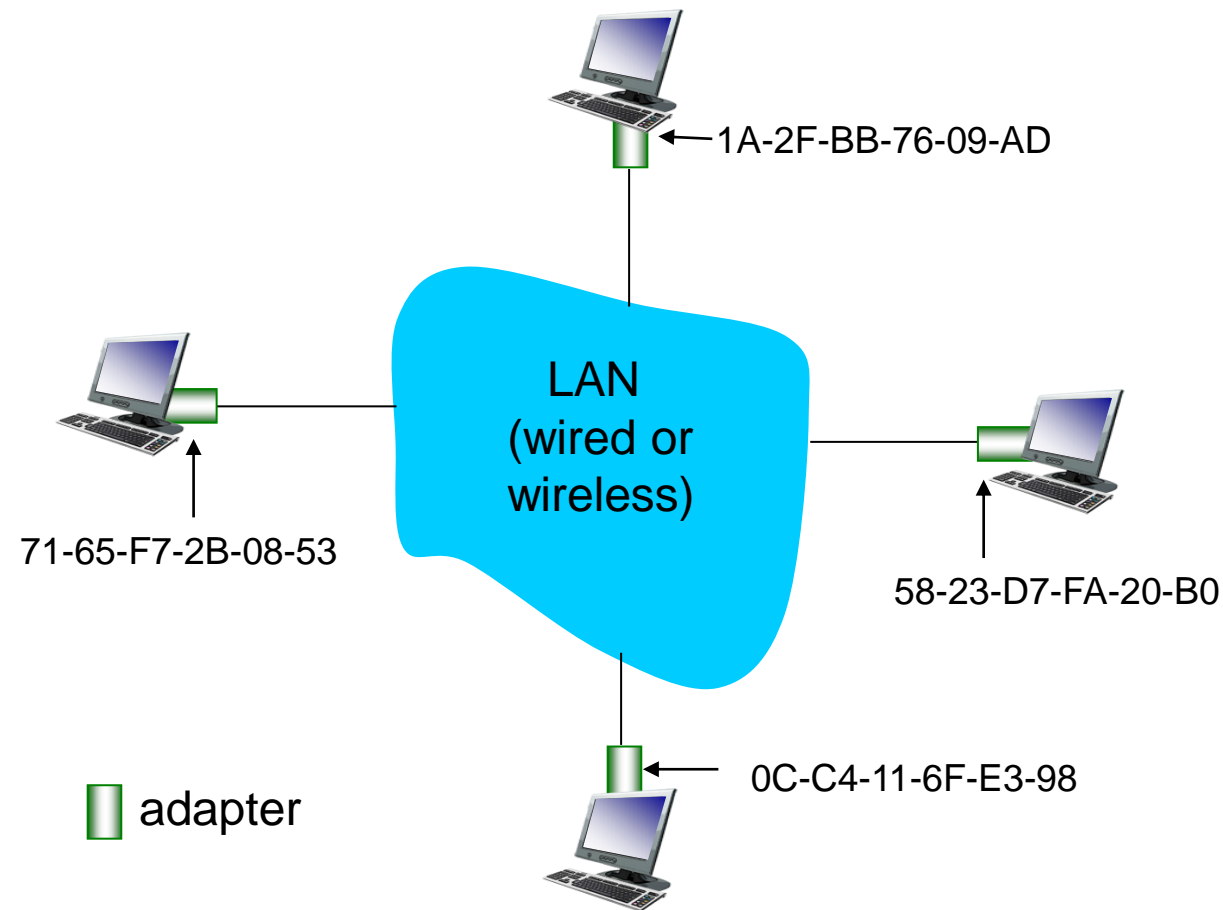
How does the sending host determine the MAC address for the destination host with its IP address ?

It uses ARP. An ARP module in the sending host takes any IP address on the same LAN as input, and returns the corresponding MAC address

ARP resolves IP addresses only for hosts and router interfaces on the same subnet.

LAN addresses and ARP

each adapter on LAN has
unique *LAN / MAC* address



- IEEE manages the MAC address space.
- In particular, when a company wants to manufacture adapters, it purchases a chunk of the address space consisting of 224 addresses for a nominal fee.
- IEEE allocates the chunk of 224 addresses by fixing the first 24 bits of a MAC address and letting the company create unique combinations of the last 24 bits for each adapter.
- An adapter's MAC address has a flat structure(not hierarchical)
- 6 byte (48 bit) hexadecimal notation
- doesn't change no matter where the adapter goes

When an adapter wants to send a frame to some destination adapter, the sending adapter inserts the destination adapter's MAC address into the frame and then sends the frame into the LAN

Because there are both network-layer addresses (for example, Internet IP addresses) and link-layer addresses (that is, MAC addresses), there is a need to translate between them. For the Internet, this is the job of the Address Resolution Protocol (ARP)

LAN addresses (more)

- ❖ MAC address allocation administered by IEEE
- ❖ manufacturer buys portion of MAC address space (to assure uniqueness)
- ❖ analogy:
 - MAC address: like Social Security Number
 - IP address: like postal address
- ❖ MAC flat address → portability
 - can move LAN card from one LAN to another
- ❖ IP hierarchical address *not* portable
 - address depends on IP subnet to which node is attached


ARP: address resolution protocol

ARP table: each IP node (host, router) on LAN has table

- IP/MAC address mappings for some LAN nodes:

< IP address; MAC address; TTL >

- TTL (Time To Live): time after which address mapping will be forgotten (typically 20 min)



Question: how to determine interface's MAC address, knowing its IP address?

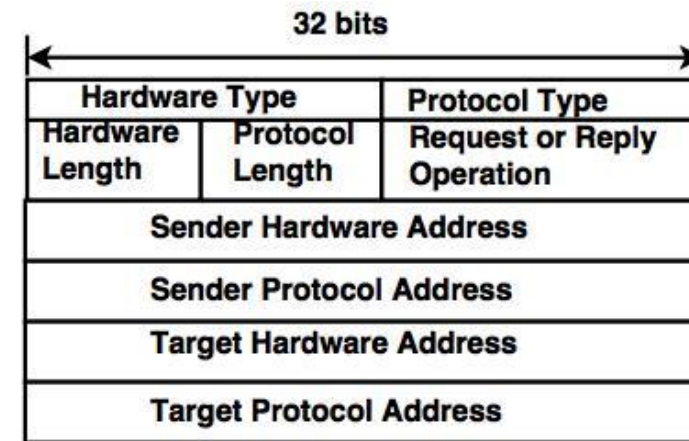
ARP protocol: same LAN

- A wants to send datagram to B
 - B's MAC address not in A's ARP table.
- A **broadcasts** ARP query packet, **containing B's IP address**
 - dest MAC address = FF-FF-FF-FF-FF-FF
 - all nodes on LAN receive ARP query
- B receives ARP packet, **replies to A with its (B's) MAC address**
 - frame sent to A's MAC address (unicast)

- A caches (saves) IP-to-MAC address pair in its **ARP table** until information becomes old (times out)
 - soft state: information that times out (goes away) unless refreshed
- ARP is “plug-and-play”:
 - nodes create their ARP tables *without intervention from net administrator*

IP Address	MAC Address	TTL
222.222.222.221	88-B2-2F-54-1A-0F	13:45:00
222.222.222.223	5C-66-AB-90-75-B1	13:52:00

Figure 5.18 ♦ A possible ARP table in 222.222.222.220




ARP Packet

Addressing: routing to another LAN


walkthrough: send datagram from A to B via R

- focus on addressing – at IP (datagram) and MAC layer (frame)
- assume A knows B's IP address
- assume A knows IP address of first hop router, R
- assume A knows R's MAC address




Addressing: routing to another LAN

- ❖ A creates IP datagram with IP source A, destination B
- ❖ A creates link-layer frame with R's MAC address as dest, frame contains A-to-B IP datagram




Addressing: routing to another LAN

- ❖ frame sent from A to R
- ❖ frame received at R, datagram removed, passed up to IP




Addressing: routing to another LAN

- ❖ R forwards datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram




Addressing: routing to another LAN

- ❖ R forwards datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram



Addressing: routing to another LAN

- ❖ R forwards datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram



Link layer, LANs: outline

5.1 introduction, services

5.2 error detection, correction

5.3 multiple access protocols

5.4 LANs

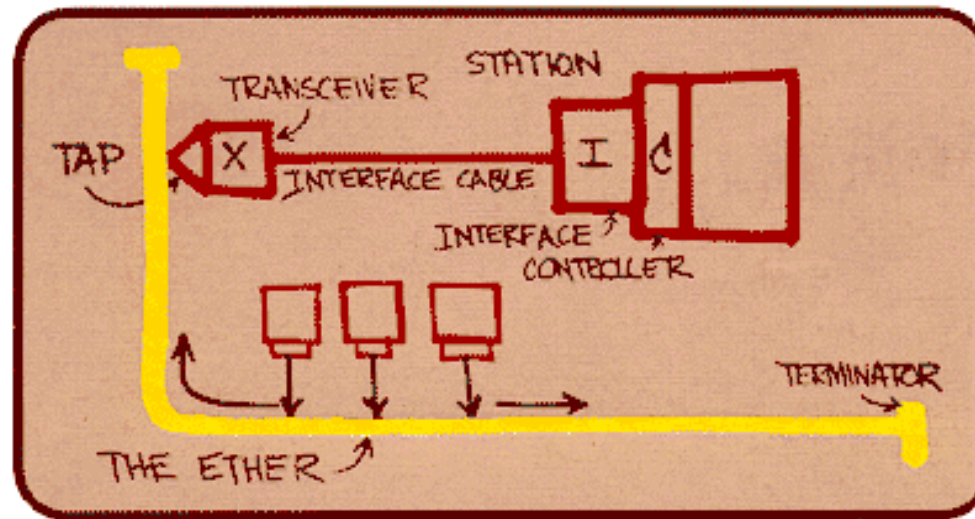
- addressing, ARP
- Ethernet
- switches
- VLANs

Ethernet

“dominant” wired LAN technology:

- cheap \$20 for NIC
- first widely used LAN technology
- simpler, cheaper than token LANs and ATM
- kept up with speed race: 10 Mbps – 10 Gbps


- Ethernet is by far the most prevalent wired LAN technology, and it is likely to remain so for the foreseeable future
- Ethernet has been to local area networking what the Internet has been to global Networking
- Ethernet was the first widely deployed high-speed LAN




Metcalfe's Ethernet sketch

Ethernet: physical topology

- *bus*: popular through mid 90s
 - all nodes in **same collision domain** (can collide with each other)
- *star*: prevails today
 - active *switch* in center
 - each “**spoke**” runs a (separate) **Ethernet protocol** (nodes do not collide with each other)




bus: coaxial cable



Ethernet frame structure

sending adapter **encapsulates** **IP datagram** (or other network layer protocol packet) in **Ethernet frame**



preamble:

- ❖ 8 byte (7 bytes with pattern **10101010** followed by one byte with pattern **10101011**)
- The first 7 bytes of the preamble serve to “wake up” the receiving adapters and to synchronize their clocks to that of the sender’s clock- used to synchronize receiver, sender clock rates

Ethernet frame structure (more)

- ❖ *addresses*: 6 byte source, destination MAC addresses
 - if adapter receives frame with matching destination address, or with broadcast address (e.g. ARP packet), it passes data in frame to network layer protocol
 - otherwise, adapter discards frame
- ❖ *type*: indicates higher layer protocol (mostly IP but others possible, e.g., Novell IPX, AppleTalk)
- ❖ *CRC*: cyclic redundancy check at receiver
 - error detected: frame is dropped



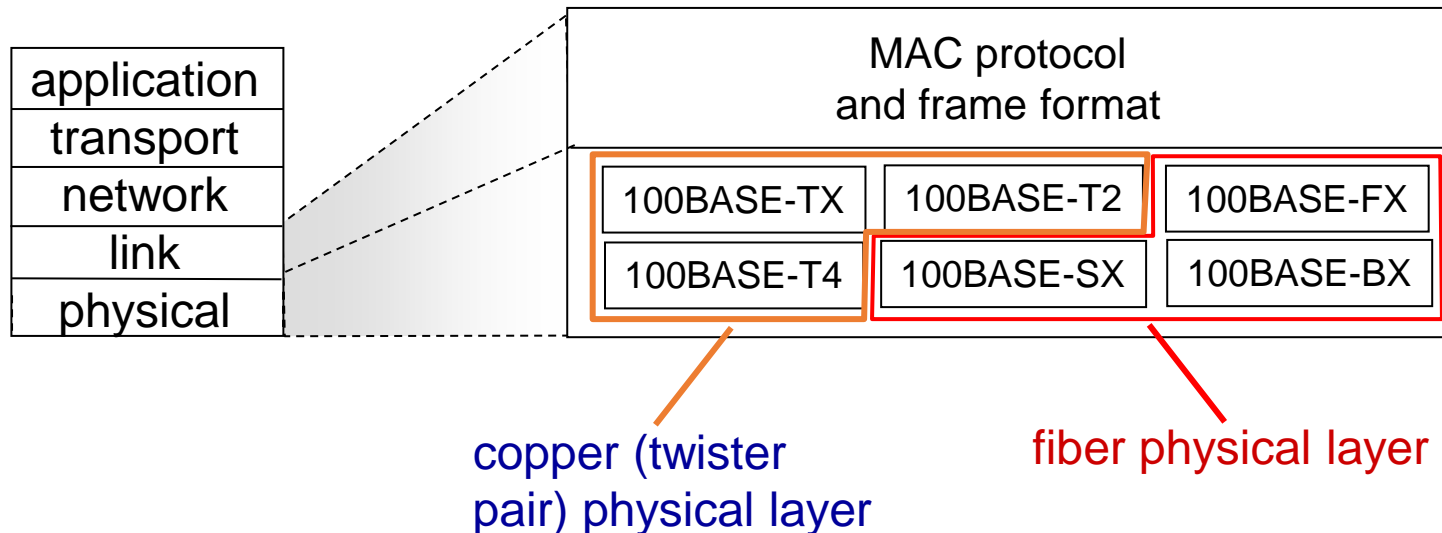
Ethernet: unreliable, connectionless

- *connectionless*: no handshaking between sending and receiving NICs
- *unreliable*: receiving NIC doesn't send acks or nacks to sending NIC
 - data in dropped frames recovered only if initial sender uses higher layer rdt (e.g., TCP), otherwise dropped data lost
- Ethernet's MAC protocol: unslotted *CSMA/CD with binary backoff*

802.3 Ethernet standards: link & physical layers

❖ *many* different Ethernet standards

- common MAC protocol and frame format
- different speeds: 2 Mbps, 10 Mbps, 100 Mbps, 1Gbps, 10G bps
- different physical layer media: fiber, cable



- The first part of the acronym refers to the **speed of the standard: 10, 100, 1000, or 10G, for 10 Megabit** (per second), 100 Megabit, Gigabit, and 10 Gigabit Ethernet, respectively.
- “**BASE**” refers to **baseband Ethernet**, meaning that the physical media only carries Ethernet traffic; almost all of the 802.3 standards are for baseband Ethernet.
- The final part of the acronym refers to the **physical media itself**; Ethernet is both a link-layer and a physical-layer specification and is carried over a variety of physical media including coaxial cable, copper wire, and fiber. Generally, a “T” refers to twisted-pair copper wires.

Link layer, LANs: outline

5.1 introduction, services

5.2 error detection, correction

5.3 multiple access protocols

5.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

5.5 link virtualization:
MPLS

5.6 data center
networking


5.7 a day in the life of a
web request

Ethernet switch

- link-layer device: takes an *active* role
 - store, forward Ethernet frames
 - examine incoming frame's MAC address, *selectively* forward frame to one-or-more outgoing links when frame is to be forwarded on segment, uses CSMA/CD to access segment
- *transparent*
 - hosts are unaware of presence of switches
- *plug-and-play, self-learning*
 - switches do not need to be configured

Switch: *multiple* simultaneous transmissions

- hosts have **dedicated, direct** connection to switch
- switches **buffer** packets
- **Ethernet protocol** used on *each* incoming link, but no collisions; full duplex
 - **each link is its own collision domain**
- **switching**: A-to-A' and B-to-B' can transmit simultaneously, without collisions



switch with six interfaces
(1,2,3,4,5,6)


Switch forwarding table

Q: how does switch know A' reachable via interface 4, B' reachable via interface 5?

- ❖ A: each switch has a **switch table**, each entry:
 - (MAC address of host, interface to reach host, time stamp)
 - looks like a routing table!

Q: how are entries created, maintained in switch table?


- something like a routing protocol?



switch with six interfaces
(1,2,3,4,5,6)

Forwarding Table

MAC Address	Port
00:10:5A:45:15:B5	1
00:10:5A:45:27:DF	2
00:10:5A:45:65:A1	3
00:10:5A:45:43:A6	4
00:10:5A:45:10:CC	5
00:10:5A:45:11:CA	6

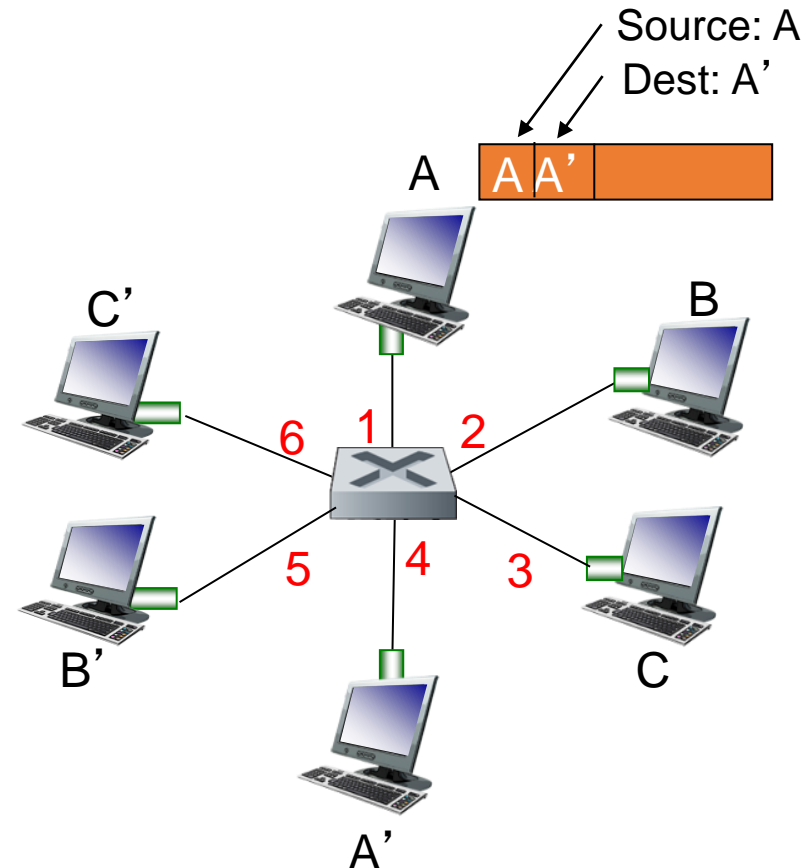


switch filtering and forwarding work

- To understand how switch filtering and forwarding work, suppose a frame with destination address DD-DD-DD-DD-DD-DD arrives at the switch on interface x . The switch indexes its table with the MAC address DD-DD-DD-DD-DD-DD. There are three possible cases:
 - **There is no entry in the table for DD-DD-DD-DD-DD-DD. In this case, the switch forwards copies of the frame to the output buffers preceding *all* interfaces except for interface x . In other words, if there is no entry for the destination address, the switch broadcasts the frame.**
 - **There is an entry in the table, associating DD-DD-DD-DD-DD-DD with interface x . In this case, the frame is coming from a LAN segment that contains adapter DD-DD-DD-DD-DD-DD. There being no need to forward the frame to any of the other interfaces, the switch performs the filtering function by discarding the frame.**
 - **There is an entry in the table, associating DD-DD-DD-DD-DD-DD with interface y . In this case, the frame needs to be forwarded to the LAN segment attached to interface y . The switch performs its forwarding function by putting the frame in an output buffer that precedes interface y**

Switch: self-learning

- switch *learns* which hosts can be reached through which interfaces
 - when frame received, switch “*learns*” location of sender: incoming LAN segment
- *records* sender/location pair in switch table



MAC addr	interface	TTL
A	1	60

*Switch table
(initially empty)*

self-learning

- Switch table is built automatically, dynamically, and autonomously—without any intervention from a network administrator or from a configuration protocol. In other words, switches are **self-learning**
 1. The switch table is initially empty.
 2. For each incoming frame received on an interface, the switch stores in its table (1) the MAC address in the frame's *source address field*, (2) the interface from which the frame arrived, and (3) the current time. In this manner the switch records in its table the LAN segment on which the sender resides. If every host in the LAN eventually sends a frame, then every host will eventually get recorded in the table.
 3. The switch deletes an address in the table if no frames are received with that address as the source address after some period of time (the **aging time**). In this manner, if a PC is replaced by another PC (with a different adapter), the MAC address of the original PC will eventually be purged from the switch table.

Switch: frame filtering/forwarding

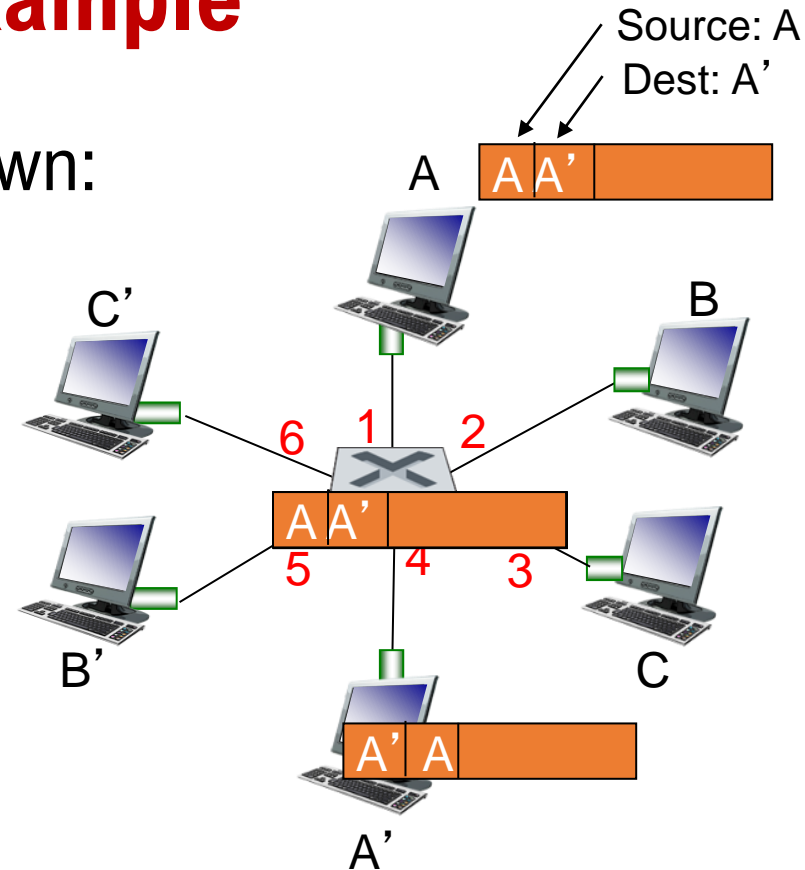
when frame received at switch:

1. **record** incoming link, MAC address of sending host
2. **index** switch table using MAC destination address
3. if entry found for **destination**
 then {
 if destination on segment from which frame arrived
 then drop frame
 else forward frame on interface indicated by entry
 }
 else flood /* **forward on all interfaces except arriving interface */**

Self-learning, forwarding: example

- frame destination, A', location unknown:

- flood*
- ❖ destination A location known: *selectively send on just one link*




MAC addr	interface	TTL
A	1	60
A'	4	60

*switch table
(initially empty)*


Interconnecting switches

❖ switches can be connected together




Q: sending from A to G - how does S_1 know to forward frame destined to G via S_4 and S_3 ?

❖ A: self learning! (works exactly the same as in single-switch case!)



Self-learning multi-switch example


Suppose C sends frame to I, I responds to C



- ❖ Q: show switch tables and packet forwarding in S₁, S₂, S₃, S₄



Institutional network



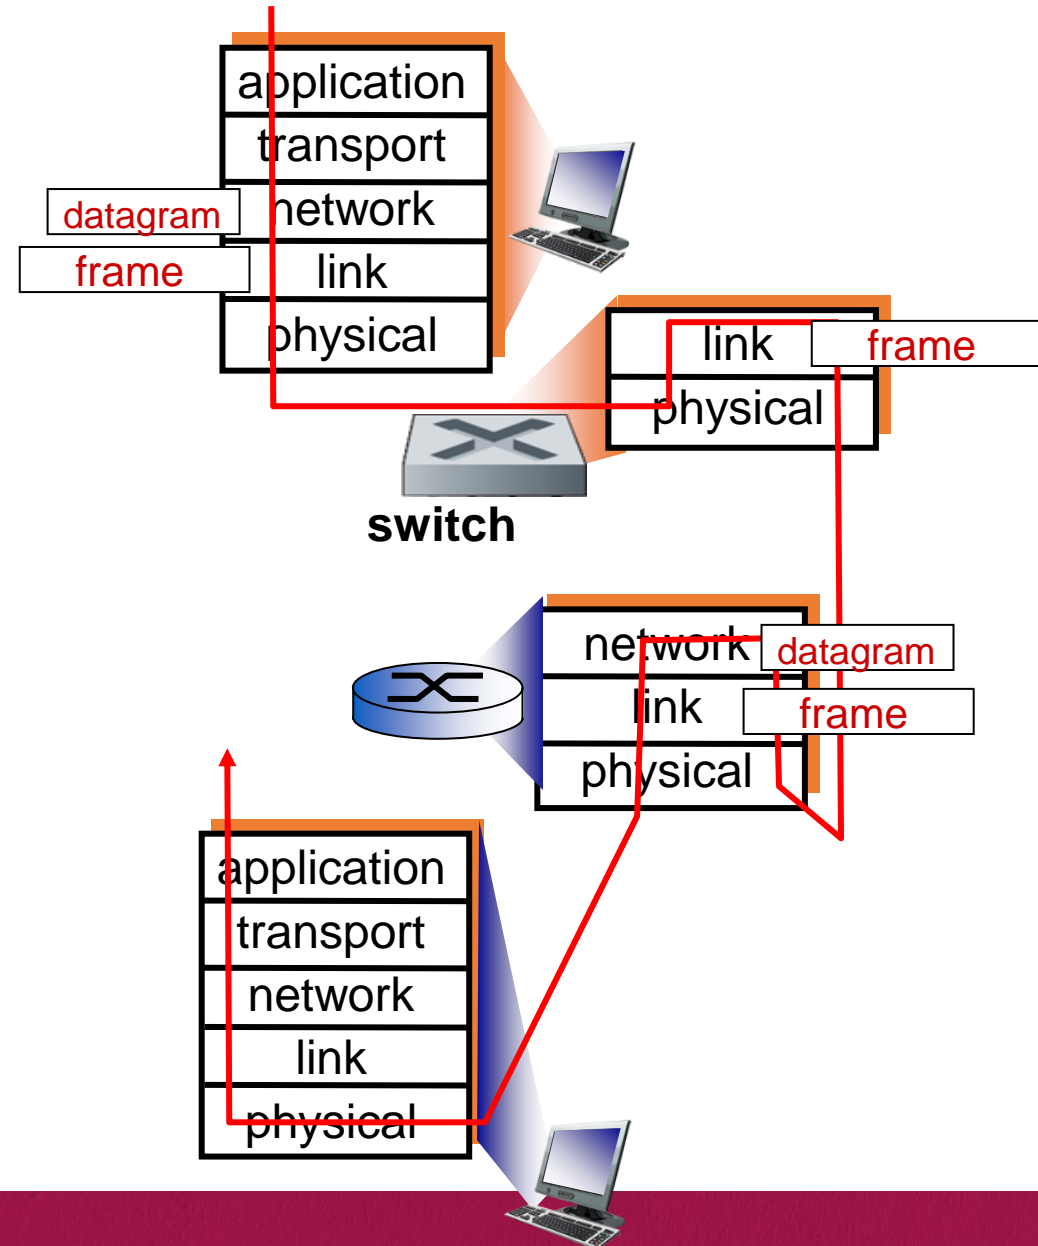
Switches vs. routers

both are store-and-forward:

- **routers:** network-layer devices (examine network-layer headers)
- **switches:** link-layer devices (examine link-layer headers)

both have forwarding tables:


- **routers:** compute tables using routing algorithms, IP addresses
- **switches:** learn forwarding table using flooding, learning, MAC addresses



VLAN (Virtual Local Area Network)

- In VLAN the computers servers and other network devices are logically connected regardless of their physical location
- Improved security
- Traffic Management
- Make a network simpler

VLANs: motivation



consider:



❖ CS user moves office to EE, but wants connect to CS switch?

❖ **single broadcast domain:**


- all layer-2 broadcast traffic (ARP, DHCP, unknown location of destination MAC address) must cross entire LAN
- security/privacy, efficiency issues


Link Layer

All the network broadcast traffic are mixed in with the other departments.




The network traffic between the departments are separated.






VLANs help with traffic management.



VLANs help with traffic management.




VLANs

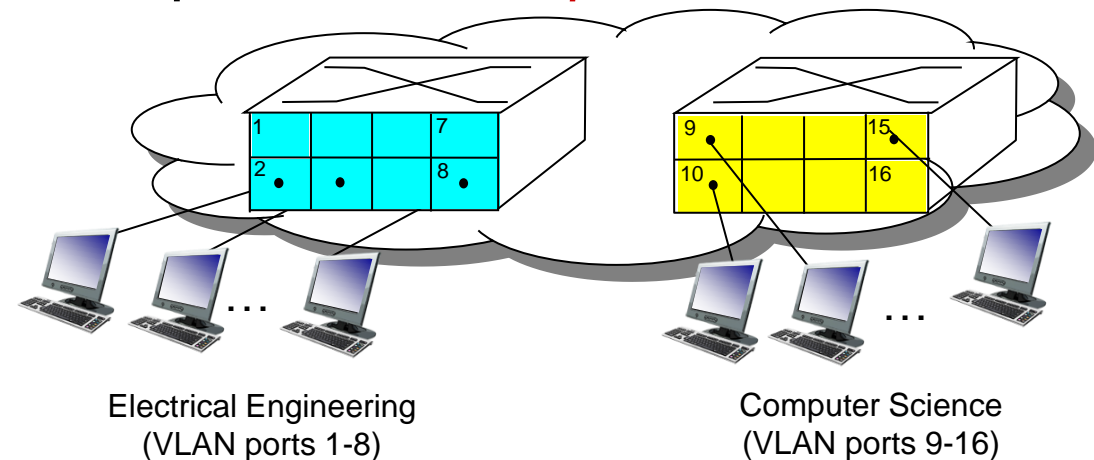
port-based VLAN: switch ports grouped (by switch management software) so that *single* physical switch can support multiple **Virtual Local Area Network**

Virtual Local Area Network

switch(es) supporting VLAN capabilities can be configured to define multiple **virtual** LANs over single physical LAN infrastructure.



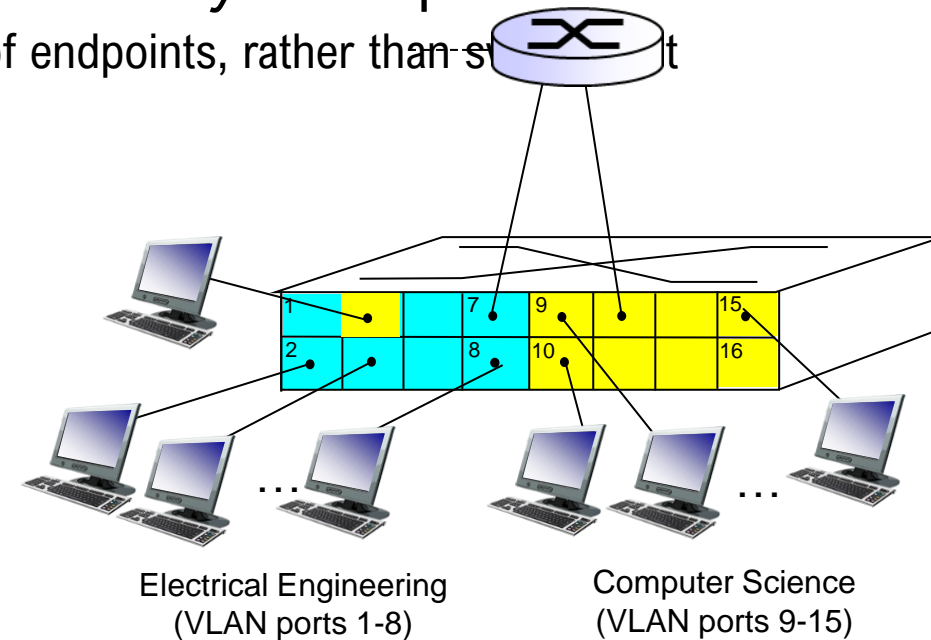
... operates as **multiple** virtual switches



Port-based VLAN

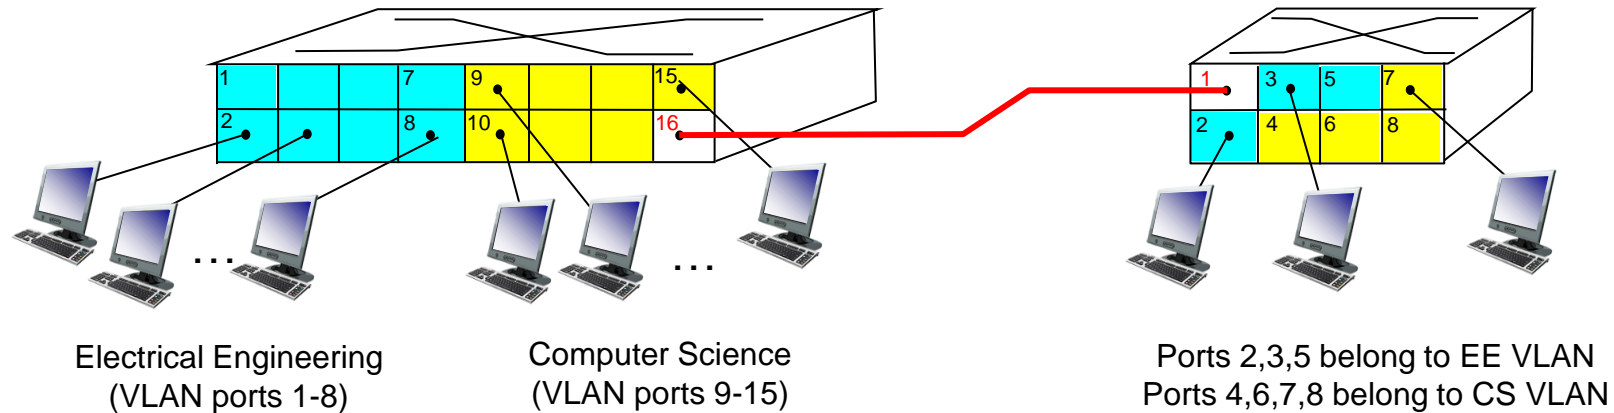
❖ *traffic isolation*: frames to/from ports 1-8 can *only* reach ports 1-8

- can also define VLAN based on MAC addresses of endpoints, rather than switch



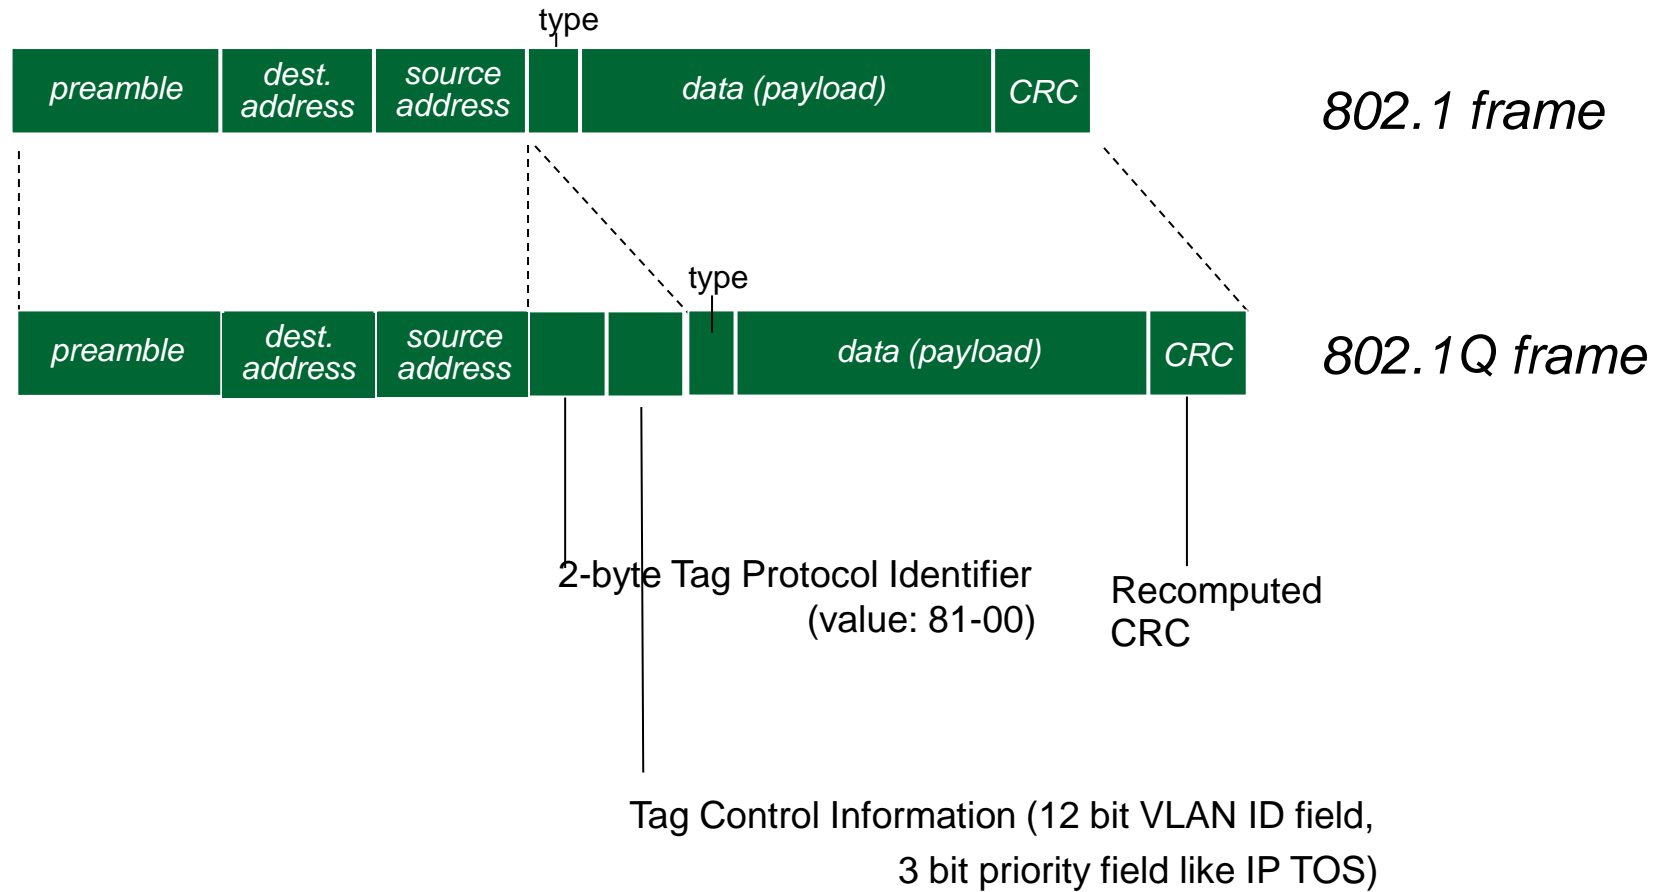
- ❖ *dynamic membership*: ports can be dynamically assigned among VLANs
- ❖ *forwarding between VLANs: done via routing* (just as with separate switches)
 - in practice vendors sell combined switches plus routers

VLANs spanning multiple switches



- **trunk port:** carries frames between VLANs defined over multiple physical switches
 - frames forwarded within VLAN between switches can't be vanilla 802.1 frames (must carry VLAN ID info)
 - 802.1Q protocol adds/removed additional header fields for frames forwarded between trunk ports

802.1Q VLAN frame format



~~Link layer, LANs: outline~~

5.1 introduction, services

5.2 error detection,
correction

5.3 multiple access
protocols

5.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

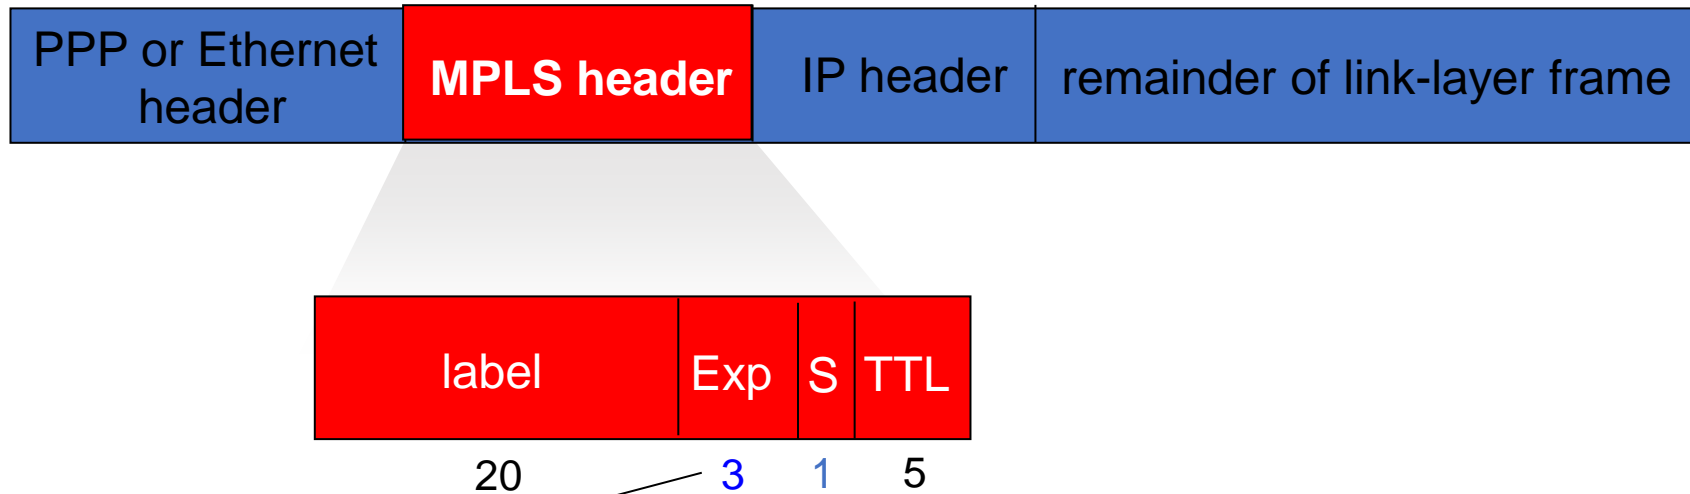
5.5 link virtualization:
MPLS

5.6 data center
networking

5.7 a day in the life of a
web request

Multiprotocol label switching (MPLS)

- ❖ initial goal: high-speed IP forwarding using **fixed length label** (instead of IP address)
 - **fast lookup** using fixed length identifier (rather than shortest prefix matching)
 - **borrowing ideas from Virtual Circuit (VC) approach**
 - but IP datagram still keeps IP address!




QoS priority (experimental) and Explicit Congestion Notification

bottom of stack flag

~~MPLS capable routers~~

- a.k.a. **label-switched router**
- forward packets to outgoing interface based only on **label value** (*don't inspect IP address*)
 - MPLS forwarding table distinct from IP forwarding tables
- ***flexibility***: MPLS forwarding decisions can *differ* from those of IP
 - use destination *and* source addresses to route flows to same destination differently (traffic engineering)
 - re-route flows quickly if link fails: pre-computed backup paths (useful for VoIP)


MPLS versus IP paths



- ❖ *IP routing: path to destination determined by destination address alone*



MPLS versus IP paths



- ❖ **IP routing:** path to destination determined by *destination address alone*




- ❖ **MPLS routing:** path to destination can be based on source *and* dest. address



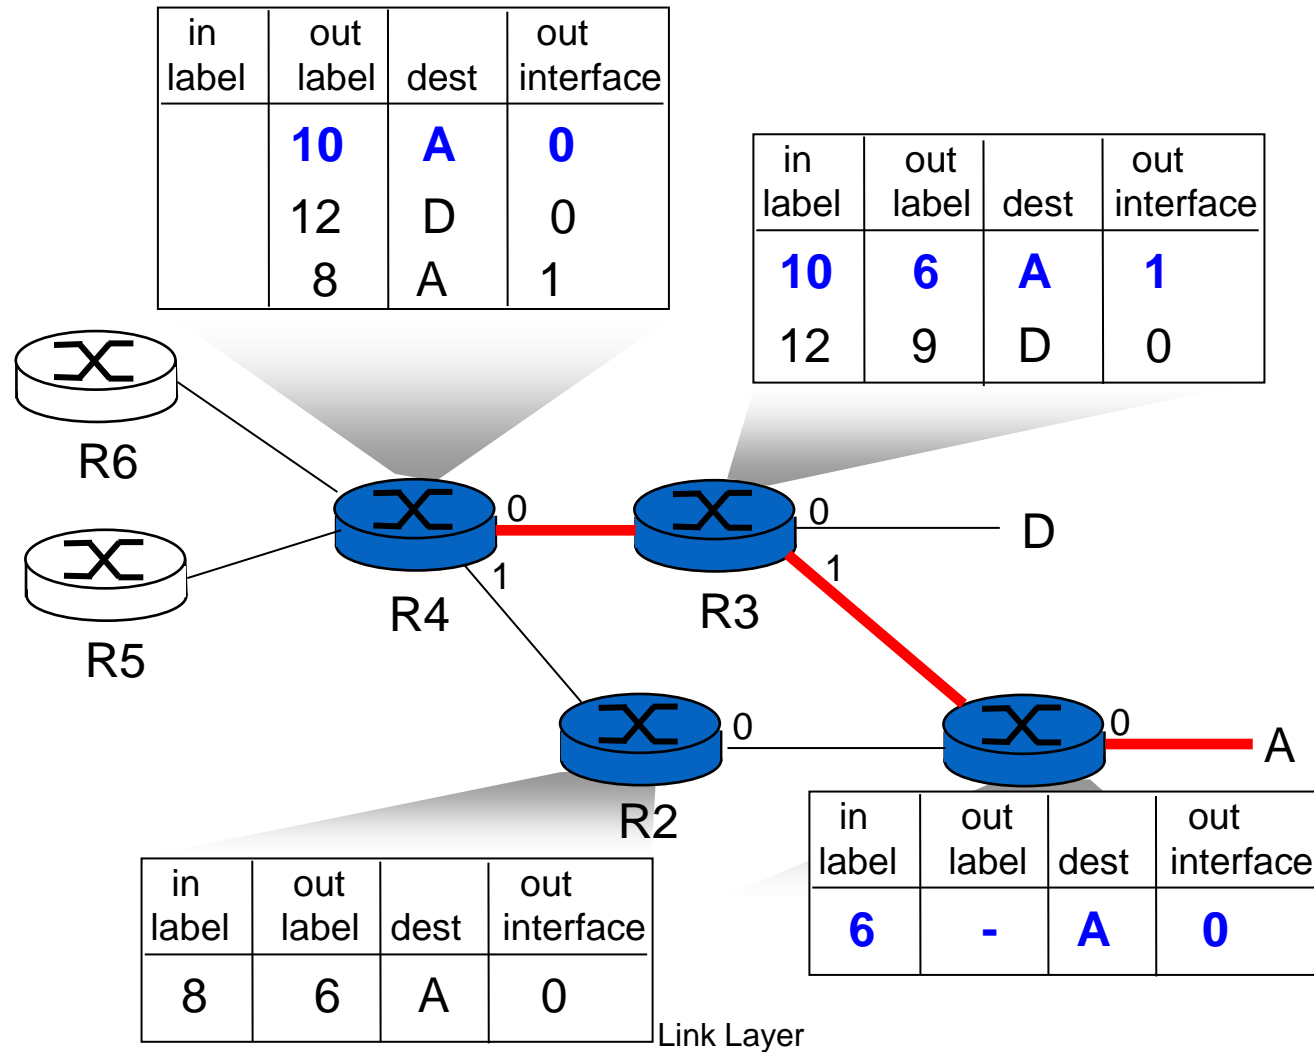
- **fast reroute:** precompute backup routes in case of link failure

MPLS ~~signaling~~

- modify OSPF, IS-IS link-state flooding protocols to carry info used by MPLS routing,
 - e.g., link bandwidth, amount of “reserved” link bandwidth
- ❖ entry MPLS router uses *RSVP-TE signaling protocol* to set up MPLS forwarding at downstream routers



MPLS forwarding tables



~~Link layer, LANs: outline~~

5.1 introduction, services

5.2 error detection,
correction

5.3 multiple access
protocols

5.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

5.5 link virtualization:
MPLS

5.6 data center
networking

5.7 a day in the life of a
web request

Data center networks

- 10's to 100's of thousands of hosts, often closely coupled, in close proximity:
 - e-business (e.g. Amazon)
 - content-servers (e.g., YouTube, Akamai, Apple, Microsoft)
 - search engines, data mining (e.g., Google)

❖ challenges:

- multiple applications, each serving massive numbers of clients
- managing/balancing load, avoiding processing, networking, data bottlenecks




Inside a 40-ft Microsoft container, Chicago data center

Data center networks


load balancer: application-layer routing

- receives external client requests
- directs workload within data center
- returns results to external client (hiding data center internals from client)



Data center networks

- ❖ rich interconnection among switches, racks:
 - increased throughput between racks (multiple routing paths possible)
 - increased reliability via redundancy



~~Link layer, LANs: outline~~

5.1 introduction, services

5.2 error detection,
correction

5.3 multiple access
protocols

5.4 LANs

- addressing, ARP
- Ethernet
- switches
- VLANs

5.5 link virtualization:
MPLS

5.6 data center
networking

5.7 a day in the life of a
web request

Synthesis: a day in the life of a web request

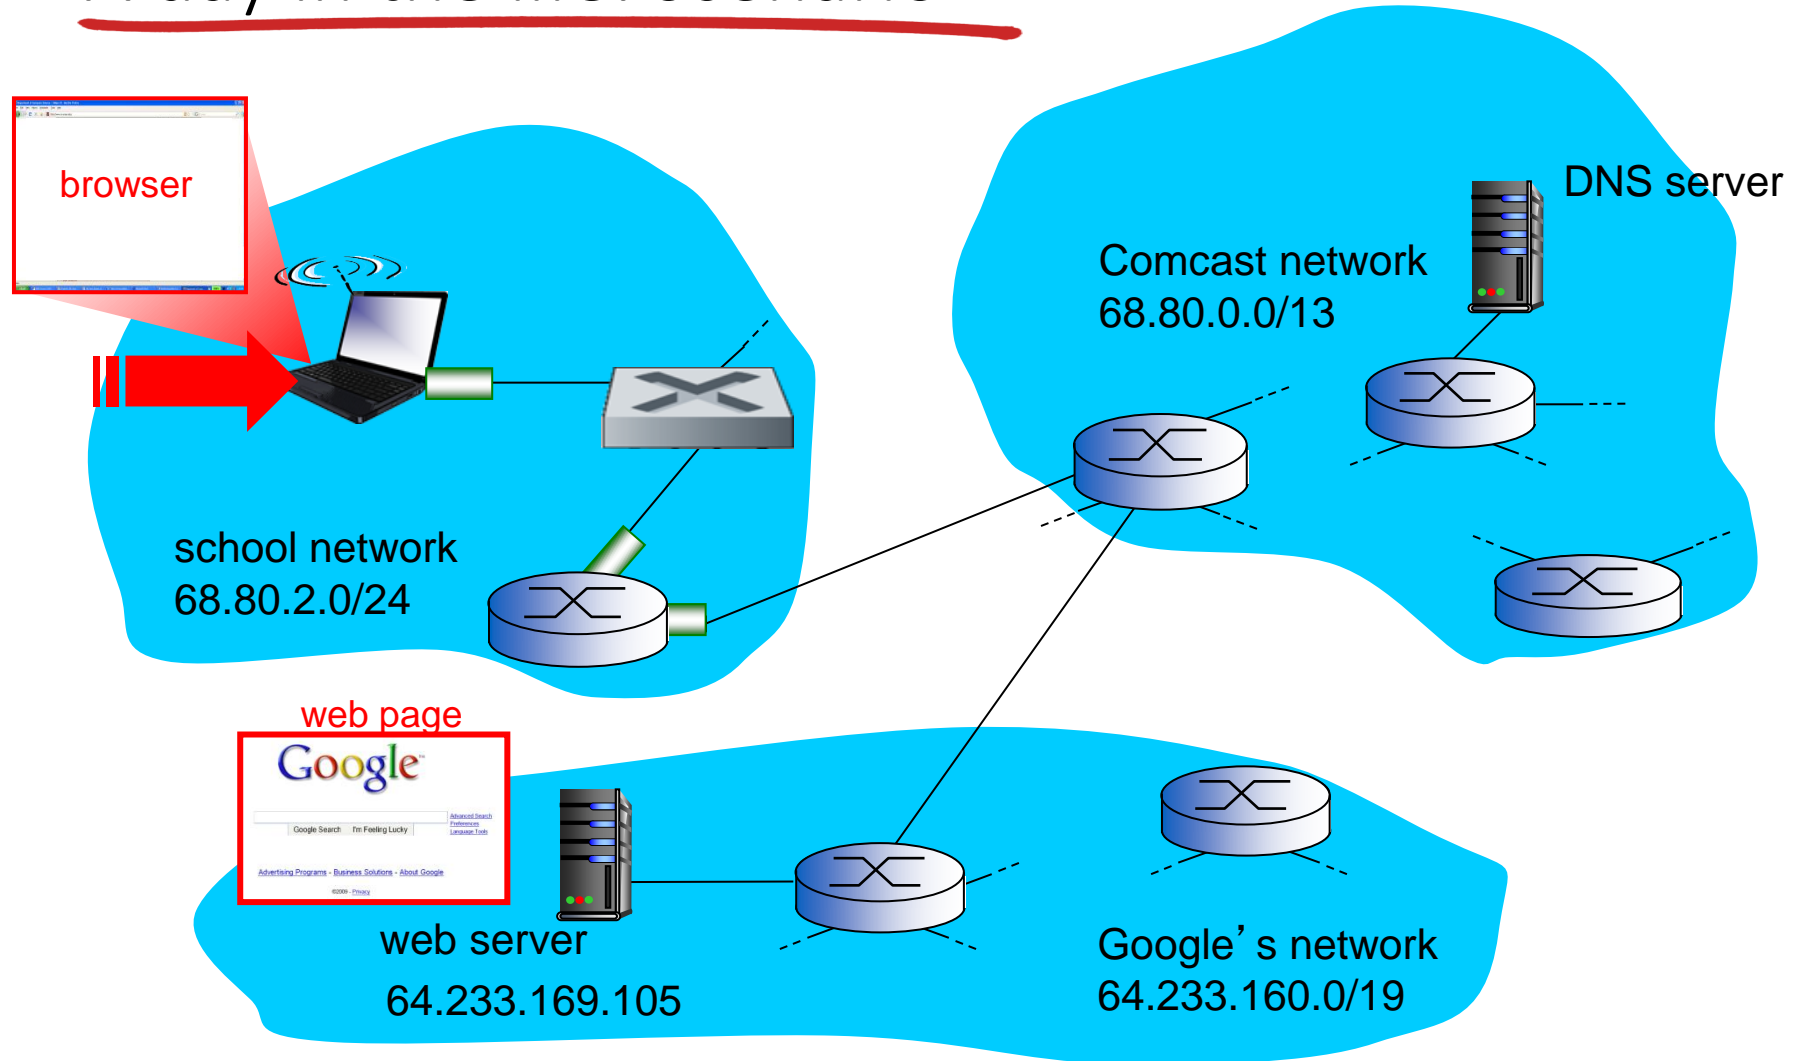
❖ journey down **protocol stack** complete!

- application, transport, network, link


❖ putting-it-all-together: synthesis!

- *goal*: identify, review, **understand protocols (at all layers)** involved in seemingly simple scenario: requesting www page
- *scenario*: student attaches laptop to campus network, requests/receives www.google.com

A day in the life: scenario




A day in the life... connecting to the Internet



- ❖ connecting laptop needs to get its own IP address, addr of first-hop router, addr of DNS server: use **DHCP**
- ❖ DHCP request **encapsulated** in **UDP**, encapsulated in **IP**, encapsulated in **802.3** Ethernet
- ❖ Ethernet frame **broadcast** (dest: FFFFFFFFFFFFFFFF) on LAN, received at router running **DHCP** server
- ❖ Ethernet **demuxed** to IP demuxed, UDP demuxed to DHCP


A day in the life... connecting to the Internet



- DHCP server formulates **DHCP ACK** containing **client's IP address**, IP address of first-hop router for client, name & IP address of DNS server
- ❖ encapsulation at DHCP server, frame forwarded (**switch learning**) through LAN, demultiplexing at client
- ❖ DHCP client receives DHCP ACK reply

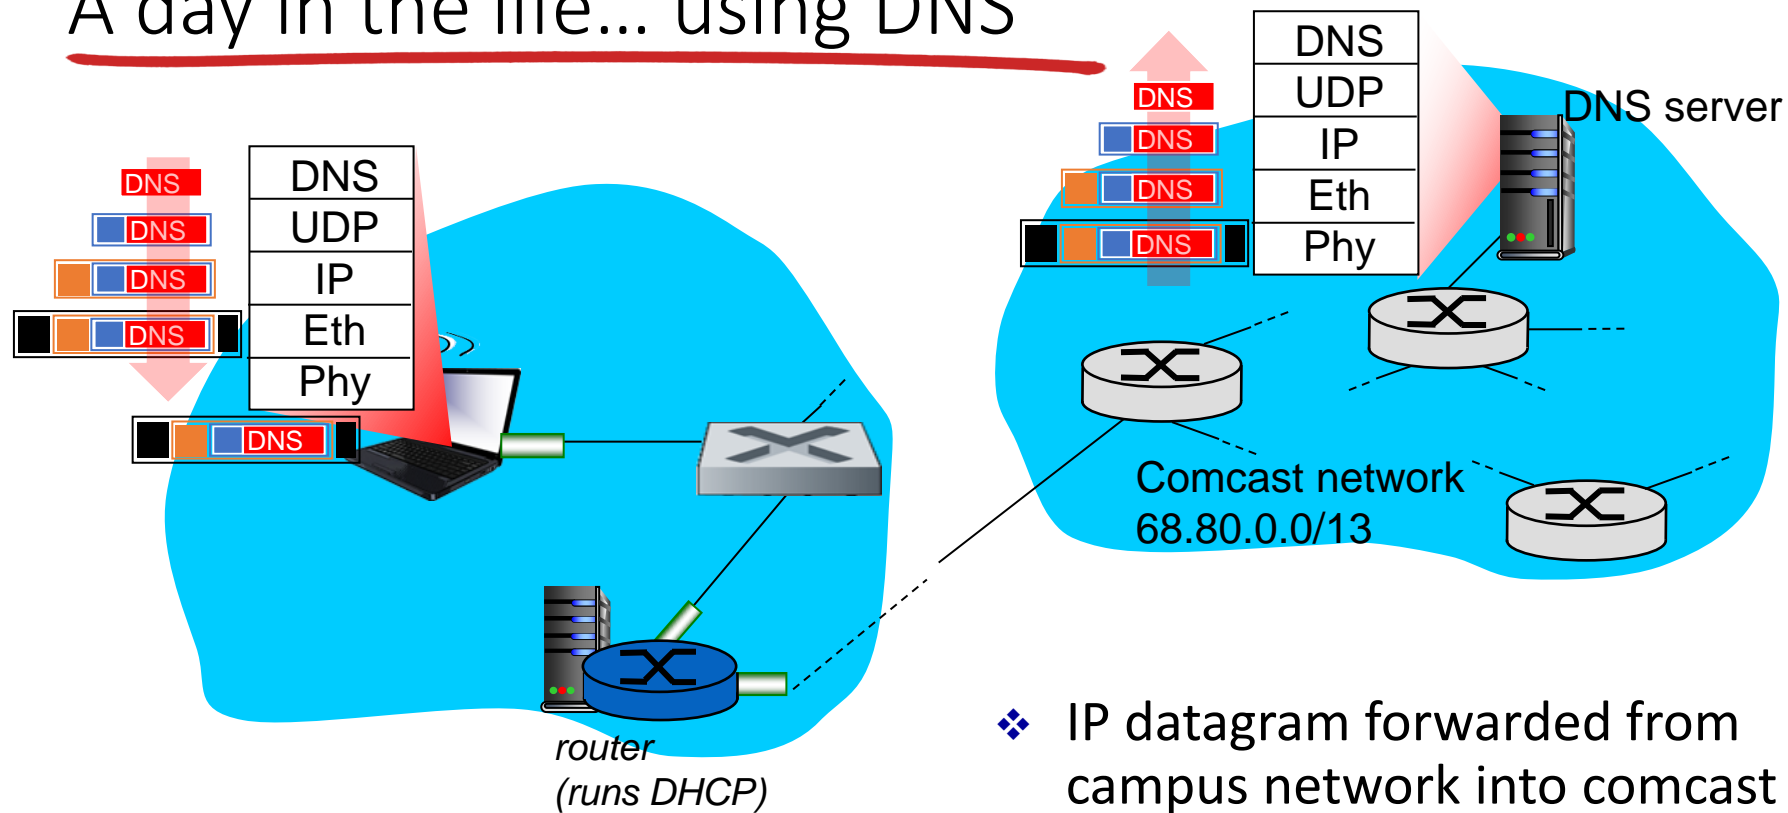
Client now has IP address, knows name & addr of DNS server, IP address of its first-hop router

A day in the life... ARP (before DNS, before HTTP)



- ❖ before sending *HTTP* request, need IP address of *www.google.com*: *DNS*
- ❖ DNS query created, encapsulated in UDP, encapsulated in IP, encapsulated in Eth. To send frame to router, need MAC address of router interface: *ARP*
- ❖ *ARP query* broadcast, received by router, which replies with *ARP reply* giving MAC address of router interface
- ❖ client now knows MAC address of first hop router, so can now send frame containing DNS query

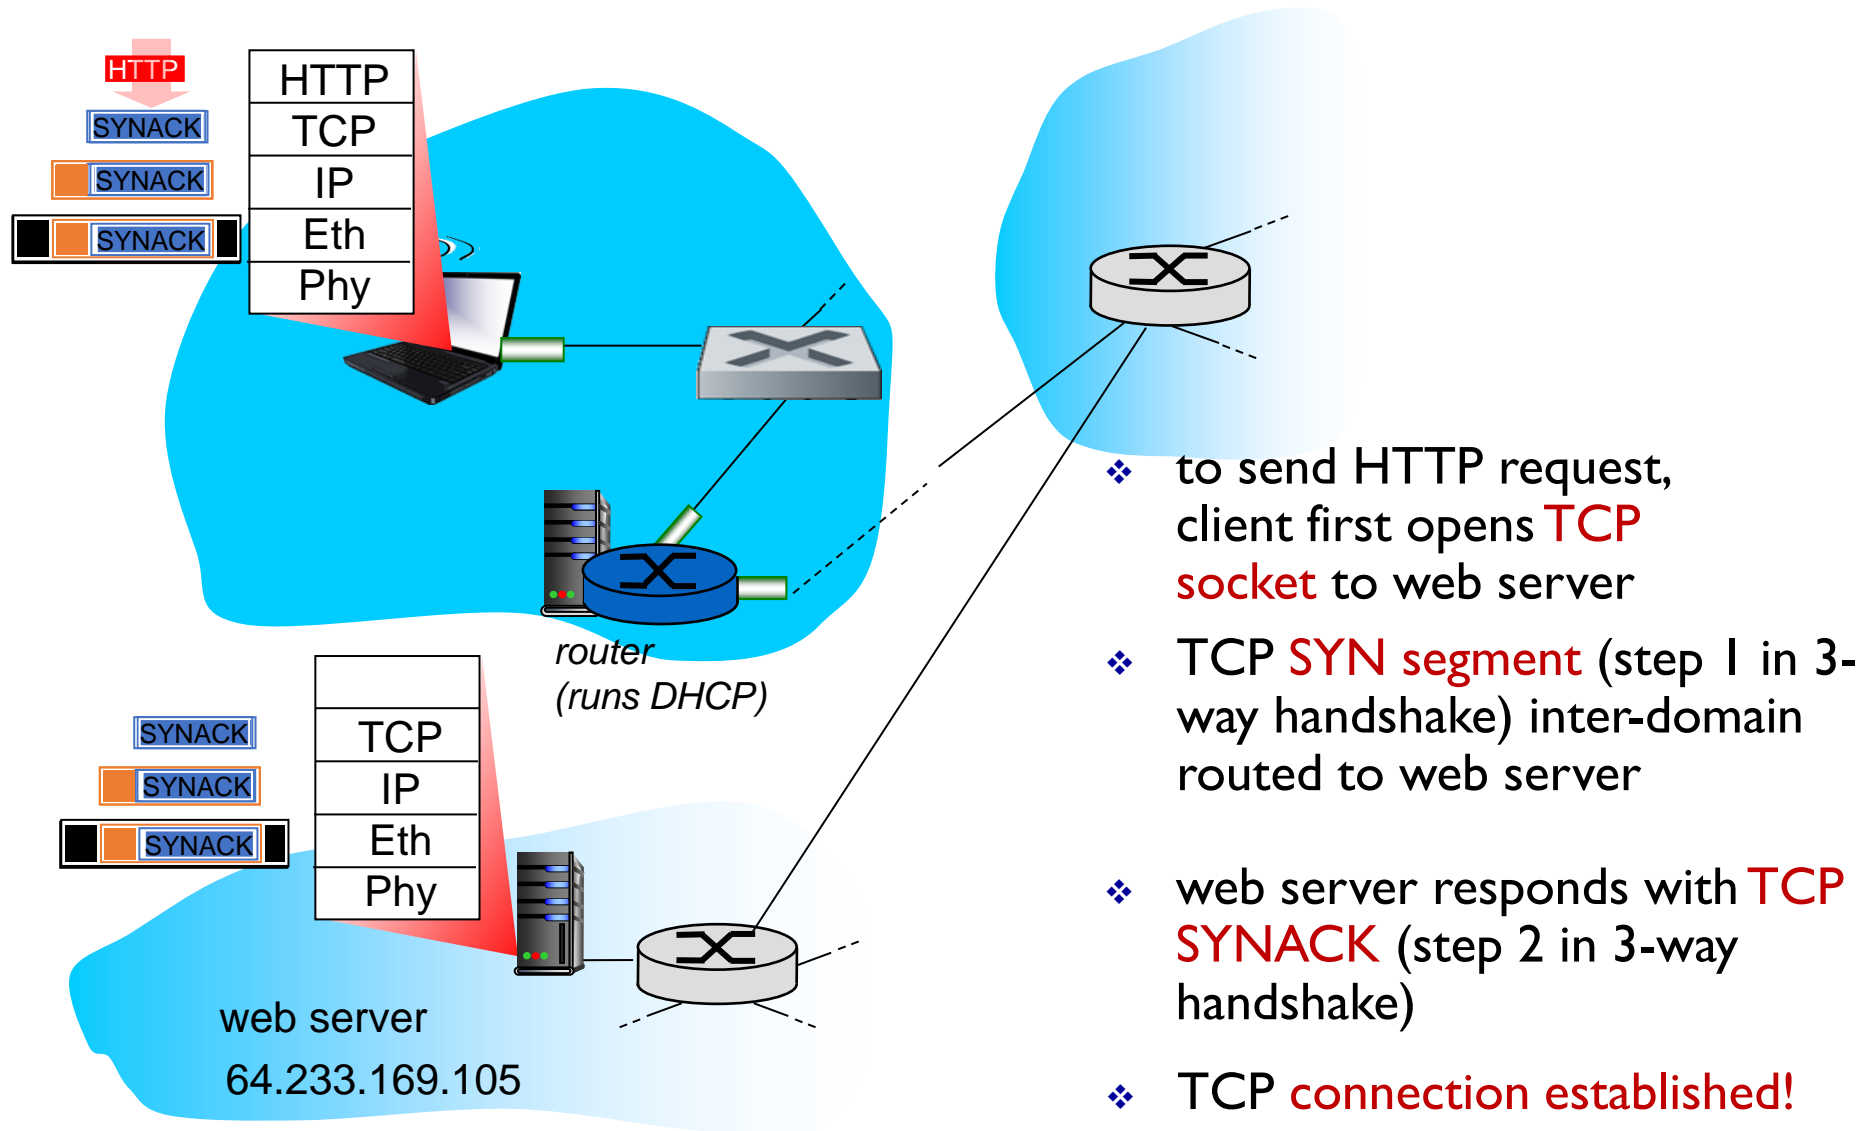
A day in the life... using DNS



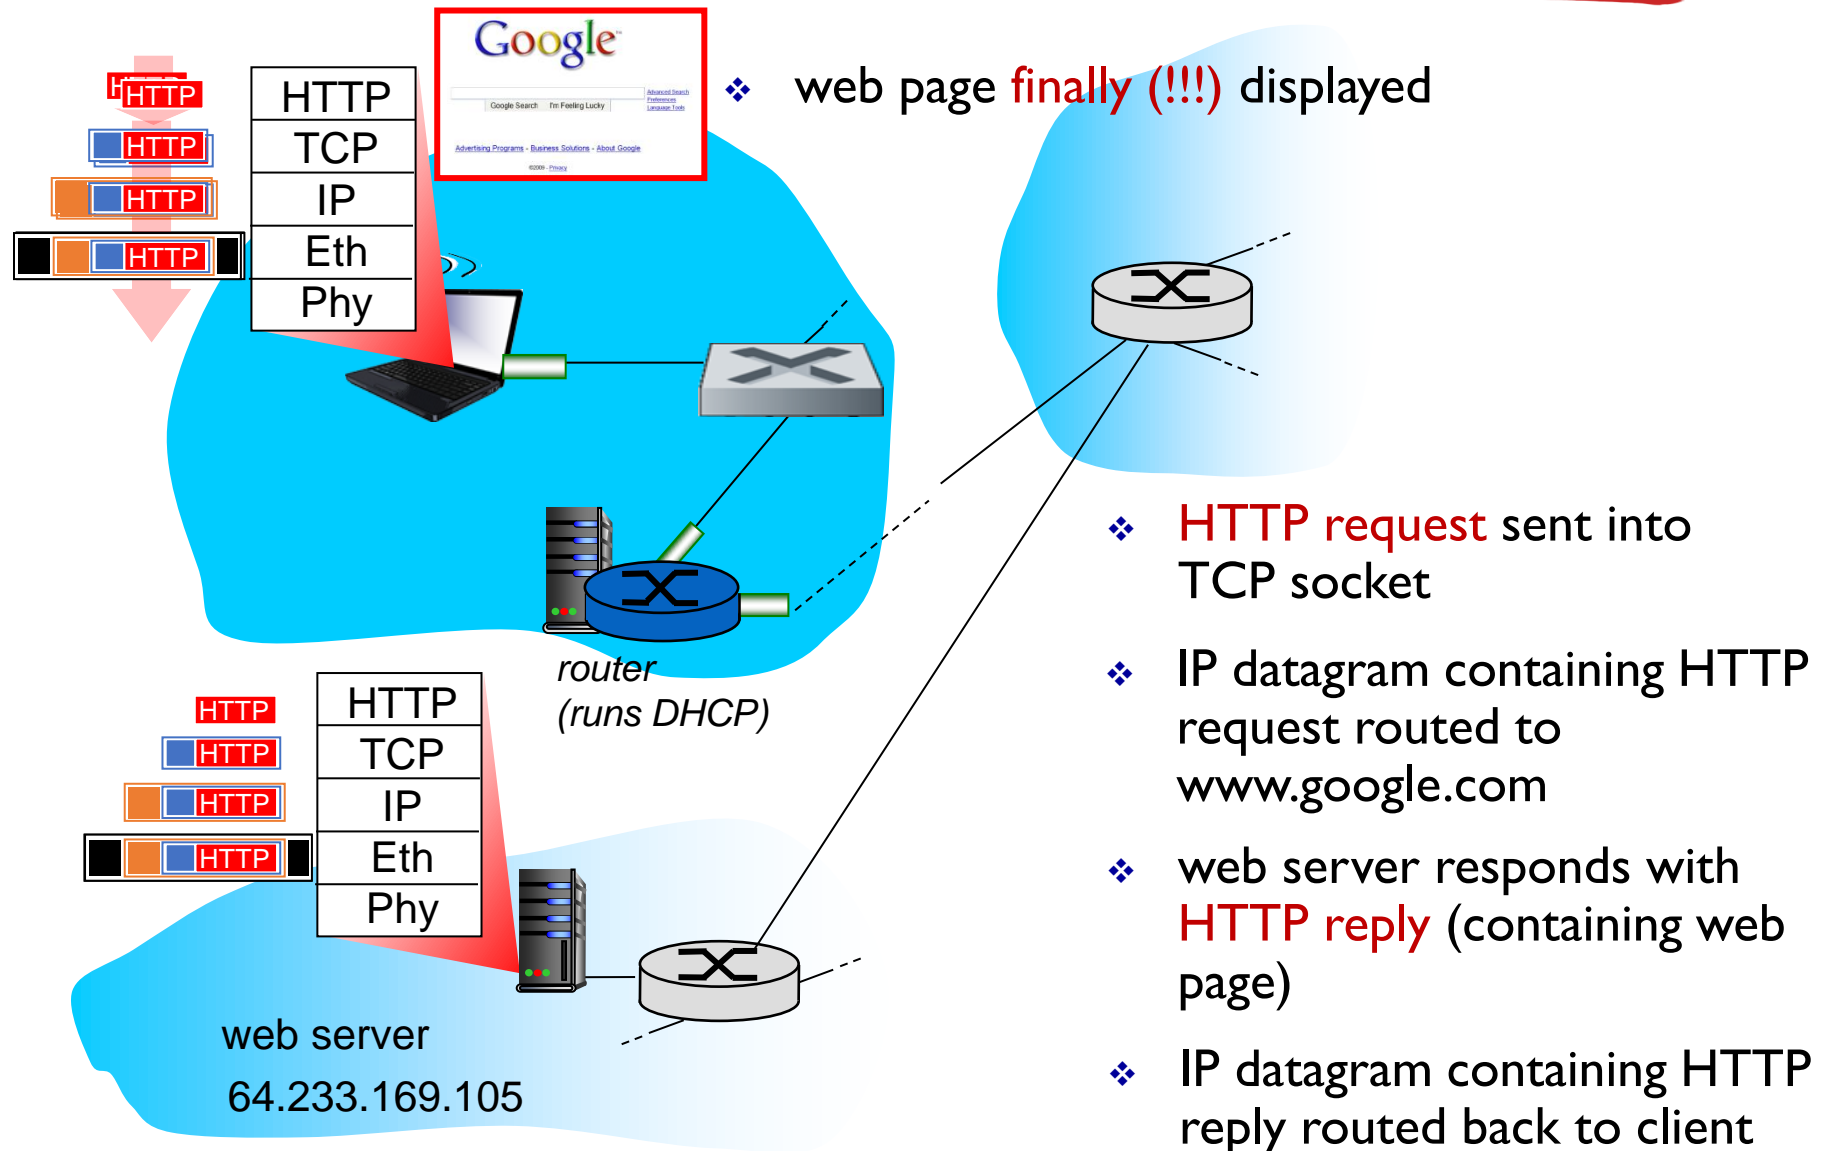
- ❖ IP datagram containing DNS query forwarded via LAN switch from client to 1st hop router

- ❖ IP datagram forwarded from campus network into comcast network, routed (tables created by **RIP**, **OSPF**, **IS-IS** and/or **BGP** routing protocols) to DNS server
- ❖ demux'ed to DNS server
- ❖ DNS server replies to client with IP address of **www.google.com**

A day in the life...TCP connection carrying HTTP



A day in the life... HTTP request/reply



~~Chapter 5: Summary~~

- ❖ principles behind data link layer services:
 - error detection, correction
 - sharing a broadcast channel: multiple access
 - link layer addressing
- ❖ instantiation and implementation of various link layer technologies
 - Ethernet
 - switched LANS, VLANs
 - virtualized networks as a link layer: MPLS
- ❖ synthesis: a day in the life of a web request

Chapter 5: let's take a breath

- journey down protocol stack *complete* (except PHY)
- solid understanding of networking principles, practice
- could stop here but *lots* of interesting topics!
 - wireless
 - multimedia
 - security
 - network management

Namah Shivaya