


Dijkstra's Algorithm

Single-Source Shortest Path Problem

Single-Source Shortest Path Problem - The problem of finding shortest paths from a source vertex v to all other vertices in the graph.



Dijkstra's algorithm

Dijkstra's algorithm - is a solution to the single-source shortest path problem in graph theory.

Works on both directed and undirected graphs. However, all edges must have nonnegative weights.

Input: Weighted graph $G=\{E,V\}$ and source vertex $v \in V$, such that all edge weights are nonnegative

Output: Lengths of shortest paths (or the shortest paths themselves) from a given source vertex $v \in V$ to all other vertices

Approach

- The algorithm computes for each vertex u the **distance** to u from the start vertex v , that is, the weight of a shortest path between v and u .
- the algorithm keeps track of the set of vertices for which the distance has been computed, called the **cloud** C
- Every vertex has a label D associated with it. For any vertex u , $D[u]$ stores an approximation of the distance between v and u . The algorithm will update a $D[u]$ value when it finds a shorter path from v to u .
- When a vertex u is added to the cloud, its label $D[u]$ is equal to the actual (final) distance between the starting vertex v and vertex u .

Dijkstra's Pseudo Code

- Graph G , weight function w , root s

DIJKSTRA(G, w, s)

1 for each $v \in V$

2 do $d[v] \leftarrow \infty$

3 $d[s] \leftarrow 0$

4 $S \leftarrow \emptyset$ \triangleright Set of discovered nodes

5 $Q \leftarrow V$

6 while $Q \neq \emptyset$

7 do $u \leftarrow \text{EXTRACT-MIN}(Q)$

8 $S \leftarrow S \cup \{u\}$

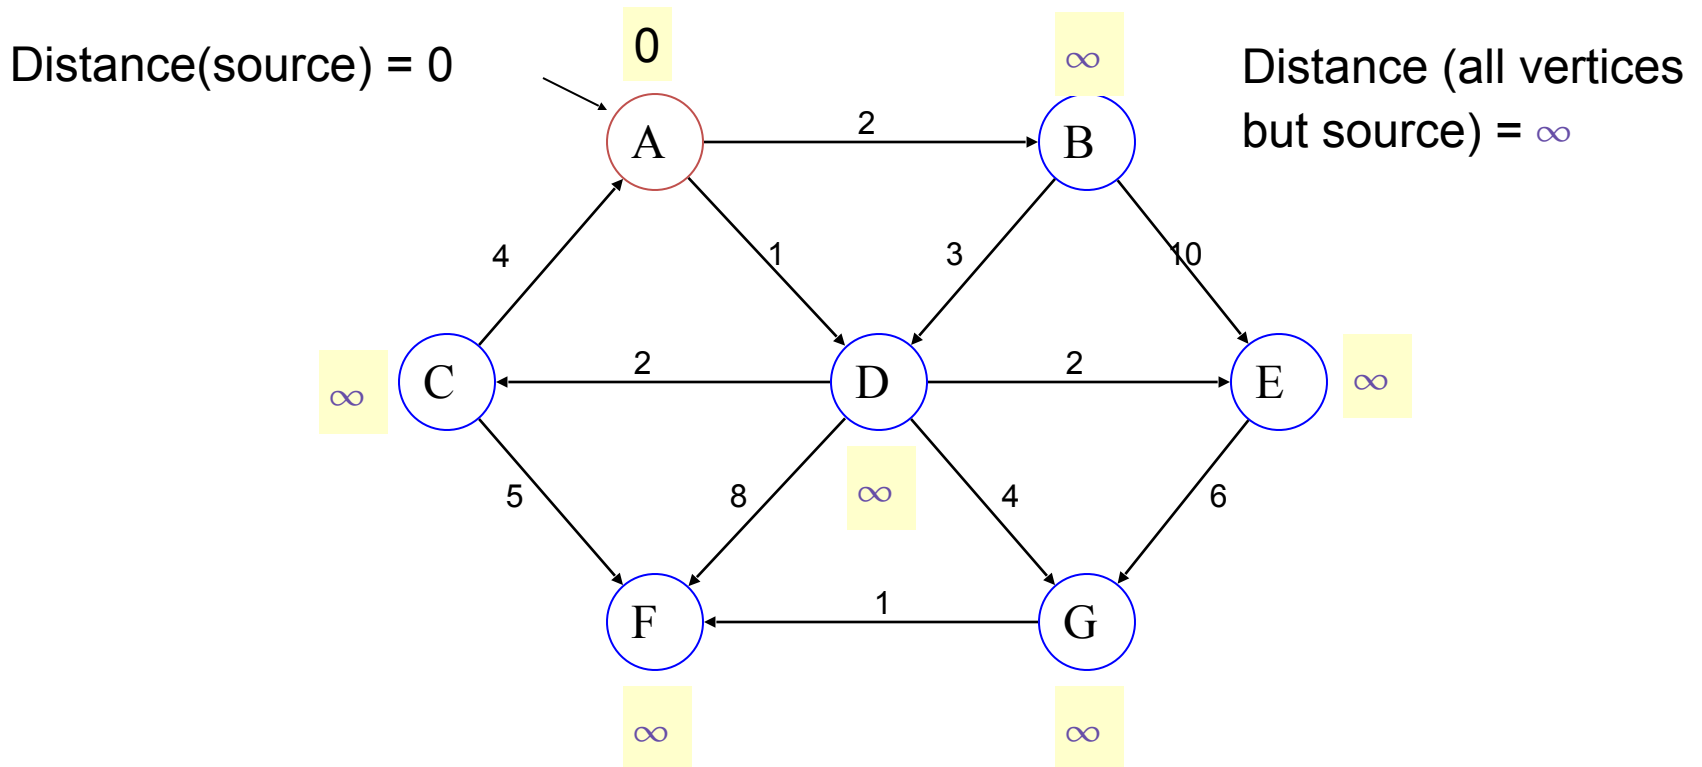
9 for each $v \in \text{Adj}[u]$

10 do if $d[v] > d[u] + w(u, v)$


11 then $d[v] \leftarrow d[u] + w(u, v)$

relaxing
edges

Example: Initialization




Example: Update neighbors' distance



Distance(B) = 2


Distance(D) = 1

Example: Remove vertex with minimum distance



Pick vertex in List with minimum distance, i.e., D

Example: Update neighbors



Distance(C) = 1 + 2 = 3


Distance(E) = 1 + 2 = 3

Distance(F) = 1 + 8 = 9

Distance(G) = 1 + 4 = 5

Example: Continued...


Pick vertex in List with minimum distance (B) and update neighbors



Note : distance(D) not updated since D is already known and distance(E) not updated since it is larger than previously computed


Example: Continued...

Pick vertex List with minimum distance (E) and update neighbors



Example: Continued...


Pick vertex List with minimum distance (C) and update neighbors



$$\text{Distance}(F) = 3 + 5 = 8$$


Example: Continued...

Pick vertex List with minimum distance (G) and update neighbors




Previous distance
↓
 $\text{Distance}(F) = \min(8, 5+1) = 6$

Example (end)




Pick vertex not in S with lowest cost (F) and update neighbors


Another Example




Another Example




Another Example




Another Example




Another Example




Another Example




Another Example



Another Example



Another Example



Is Dijkstra's algorithm correct?

- Invariant: For every vertex removed from the heap, $\text{dist}[v]$ is the actual shortest distance from s to v
 - The only time a vertex gets visited is when the distance from s to that vertex is smaller than the distance to any remaining vertex
 - Therefore, there cannot be any other path that hasn't been visited already that would result in a shorter path

Bounding the distance

- Another invariant: For each vertex v , $\text{dist}[v]$ is an upper bound on the actual shortest distance
 - start of at ∞
 - only update the value if we find a shorter distance
- An update procedure


$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$

$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$

- Can we ever go wrong applying this update rule?
 - We can apply this rule as many times as we want and will never underestimate $\text{dist}[v]$
- When will $\text{dist}[v]$ be right?
 - If u is along the shortest path to v and $\text{dist}[u]$ is correct


$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$

- $\text{dist}[v]$ will be right if u is along the shortest path to v and $\text{dist}[u]$ is correct
- Consider the shortest path from s to v




$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$

- $\text{dist}[v]$ will be right if u is along the shortest path to v and $\text{dist}[u]$ is correct
- What happens if we update all of the vertices with the above update?



$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$


- $\text{dist}[v]$ will be right if u is along the shortest path to v and $\text{dist}[u]$ is correct
- What happens if we update all of the vertices with the above update?



correct


$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$

- $\text{dist}[v]$ will be right if u is along the shortest path to v and $\text{dist}[u]$ is correct
- What happens if we update all of the vertices with the above update?




$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$

- $\text{dist}[v]$ will be right if u is along the shortest path to v and $\text{dist}[u]$ is correct
- Does the order that we update the vertices matter?




$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$

- $\text{dist}[v]$ will be right if u is along the shortest path to v and $\text{dist}[u]$ is correct
- How many times do we have to do this for vertex p_i to have the correct shortest path from s ?
 - i times




$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$

- $\text{dist}[v]$ will be right if u is along the shortest path to v and $\text{dist}[u]$ is correct
- How many times do we have to do this for vertex p_i to have the correct shortest path from s ?
 - i times




$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$

- $\text{dist}[v]$ will be right if u is along the shortest path to v and $\text{dist}[u]$ is correct
- How many times do we have to do this for vertex p_i to have the correct shortest path from s ?
 - i times




$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$

- $\text{dist}[v]$ will be right if u is along the shortest path to v and $\text{dist}[u]$ is correct
- How many times do we have to do this for vertex p_i to have the correct shortest path from s ?
 - i times




$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$

- $\text{dist}[v]$ will be right if u is along the shortest path to v and $\text{dist}[u]$ is correct
- How many times do we have to do this for vertex p_i to have the correct shortest path from s ?
 - i times



$$\text{dist}[v] = \min \{ \text{dist}[v], \text{dist}[u] + w(u, v) \}$$

- $\text{dist}[v]$ will be right if u is along the shortest path to v and $\text{dist}[u]$ is correct
- What is the longest (vertex-wise) the path from s to any node v can be?
 - $|V| - 1$ edges/vertices



Time Complexity: Using List

The simplest implementation of the Dijkstra's algorithm stores vertices in an ordinary linked list or array

- Good for dense graphs (many edges)

- $|V|$ vertices and $|E|$ edges
- Initialization $O(|V|)$
- While loop $O(|V|)$
 - Find and remove min distance vertices $O(|V|)$
- Potentially $|E|$ updates
 - Update costs $O(1)$

Total time $O(|V^2| + |E|) = O(|V^2|)$

Time Complexity: Priority Queue

For sparse graphs, (i.e. graphs with much less than $|V|^2$ edges)
Dijkstra's implemented more efficiently by *priority queue*

- Initialization $O(|V|)$ using $O(|V|)$ buildHeap
- While loop $O(|V|)$
 - Find and remove min distance vertices $O(\log |V|)$ using $O(\log |V|)$ deleteMin
- Potentially $|E|$ updates
 - Update costs $O(\log |V|)$ using decreaseKey

Total time $O(|V|\log|V| + |E|\log|V|) = O(|E|\log|V|)$

- $|V| = O(|E|)$ assuming a connected graph