

Name: Vinayak V Thayil  
Roll No:AM.EN.U4CSE21161

```
#importing pandas
import pandas as p;

#importing csv files
df=p.read_csv('diabetes.csv')

#print value of df
print(df)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..	...	...	...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
#prints last 9 values of df
df.tail(9)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
759	6	190	92	0	0	35.5	0.278	66	1
760	2	88	58	26	16	28.4	0.766	22	0
761	9	170	74	31	0	44.0	0.403	43	1
762	9	89	62	0	0	22.5	0.142	33	0
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

```
#prints first 10 values of df
df.head(9)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1

```
#prints statistical values of df
df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

```
#prints info of df
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                    768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                    768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
#returns shape of df
df.shape
```

```
(768, 9)
```

```
#returns columns of df
df.columns
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

```
#checking missing values
df.isnull()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False

```
df2=df.copy()
df2.loc[2:4, 'Pregnancies':'Glucose']=None
df2.head(7)
df2.isnull().head(7)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False
2	True	True	False	False	False	False	False	False	False
3	True	True	False	False	False	False	False	False	False
4	True	True	False	False	False	False	False	False	False
5	False	False	False	False	False	False	False	False	False
6	False	False	False	False	False	False	False	False	False

```
df2.isnull().sum()
```

```
Pregnancies      3
Glucose           3
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64
```

```
df2.isnull().sum().sum()
```

```
6
```

```
#prints values from 2 to 4
df[2:4]
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0

```
df['Outcome']
```

```
0    1
1    0
2    1
3    0
4    1
..
763  0
764  0
765  0
766  1
767  0
Name: Outcome, Length: 768, dtype: int64
```

```
#sorts BMI in descending order
df["BMI"].sort_index(ascending=False).head()
```

```
767    30.4
766    30.1
765    26.2
764    36.8
763    32.9
Name: BMI, dtype: float64
```

```
df[df.index==1]
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
1	1	85	66	29	0	26.6	0.351	31	0

```
#prints values of 1,2,3
df[df.index.isin([1,2,3])]
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0

```
#prints values of blood pressure greater than 90
df[df.BloodPressure>90][1:4]
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
10	4	110	92	0	0	37.6	0.191	30	0
24	11	143	94	33	146	36.6	0.254	51	1
29	5	117	92	0	0	34.1	0.337	38	0

```
#prints mean of df
df.mean()
```

```
Pregnancies      3.845052
Glucose          120.894531
BloodPressure     69.105469
SkinThickness    20.536458
Insulin          79.799479
BMI              31.992578
DiabetesPedigreeFunction  0.471876
Age              33.240885
Outcome          0.348958
dtype: float64
```

```
#creates new column glucose insulin ratio
df2['Glucose_insulin_ratio']=df2['Glucose']/df2['Insulin']
```

```
#prints count of each value in pregnancies
df2['Pregnancies'].value_counts()
```

```
1.0    134
0.0    110
2.0    103
3.0     75
4.0     68
5.0     57
6.0     50
7.0     45
8.0     37
9.0     28
10.0    24
11.0    11
13.0     9
12.0     9
14.0     2
15.0     1
17.0     1
Name: Pregnancies, dtype: int64
```

```
#1
#stores data in wq
wq=p.read_csv('WineQT.csv')
```

```
#prints wq
print(wq)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides \
0	7.4	0.700	0.00	1.9	0.076
1	7.8	0.880	0.00	2.6	0.098
2	7.8	0.760	0.04	2.3	0.092
3	11.2	0.280	0.56	1.9	0.075
4	7.4	0.700	0.00	1.9	0.076
...	...	...	...	...	...
1138	6.3	0.510	0.13	2.3	0.076
1139	6.8	0.620	0.08	1.9	0.068
1140	6.2	0.600	0.08	2.0	0.090
1141	5.9	0.550	0.10	2.2	0.062
1142	5.9	0.645	0.12	2.0	0.075

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
0	11.0	34.0	0.99780	3.51	0.56
1	25.0	67.0	0.99680	3.20	0.68
2	15.0	54.0	0.99700	3.26	0.65
3	17.0	60.0	0.99800	3.16	0.58
4	11.0	34.0	0.99780	3.51	0.56
...	...	...	...	...	...
1138	29.0	40.0	0.99574	3.42	0.75
1139	28.0	38.0	0.99651	3.42	0.82
1140	32.0	44.0	0.99490	3.45	0.58
1141	39.0	51.0	0.99512	3.52	0.76
1142	32.0	44.0	0.99547	3.57	0.71

	alcohol	quality	Id
0	9.4	5	0
1	9.8	5	1
2	9.8	5	2
3	9.8	6	3
4	9.4	5	4
...	...	...	...
1138	11.0	6	1592
1139	9.5	6	1593
1140	10.5	5	1594
1141	11.2	6	1595
1142	10.2	5	1597

[1143 rows x 13 columns]

```
#3
#returns shape of wq
wq.shape
```

(1143, 13)

```
#4
#prints statistical values of wq
wq.describe()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000
mean	8.311111	0.531339	0.268364	2.532152	0.086933	15.615486	45.915486
std	1.747595	0.179633	0.196686	1.355917	0.047267	10.250486	32.765486
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000
25%	7.100000	0.392500	0.090000	1.900000	0.070000	7.000000	21.000000
50%	7.900000	0.520000	0.250000	2.200000	0.079000	13.000000	37.000000
75%	9.100000	0.640000	0.420000	2.600000	0.090000	21.000000	61.000000



```
#5
#prints first 5 values of wq
wq.head(5)
```

```
fixed volatile citric residual chlorides free total
acidity acidity acid sugar chlorides sulfur sulfur
density pH sulph
dioxide dioxide
```

```
#5
#prints last 5 values of wq
wq.tail(5)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	su
1138	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	
1139	6.8	0.620	0.08	1.9	0.068	28.0	38.0	0.99651	3.42	
1140	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	
1141	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	

```
#6
#prints first 7 values of wq
wq.head(7)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulph
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	

```
#7
wq.iloc[0:7,4:7]
```

	chlorides	free sulfur dioxide	total sulfur dioxide
0	0.076	11.0	34.0
1	0.098	25.0	67.0
2	0.092	15.0	54.0
3	0.075	17.0	60.0
4	0.076	11.0	34.0
5	0.075	13.0	40.0
6	0.069	15.0	59.0

```
#8
wq.iloc[:,3:6]
```

	residual sugar	chlorides	free sulfur dioxide
0	1.9	0.076	11.0
1	2.6	0.098	25.0



#9

```
#checks for null values in wq
wq.isnull().sum()
```

```
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                 0
sulphates          0
alcohol            0
quality            0
Id                 0
dtype: int64
```

#10

```
#renames column ph to ph_value
wq.rename(columns={'pH':'potential of hydrogen'},inplace=True)
```

```
print(wq['potential of hydrogen'])
```

```
0      3.51
1      3.20
2      3.26
3      3.16
4      3.51
...
1138   3.42
1139   3.42
1140   3.45
1141   3.52
1142   3.57
Name: potential of hydrogen, Length: 1143, dtype: float64
```

#11

```
wq['total_free_ratio'] = wq['total sulfur dioxide'] / wq['free sulfur dioxide']
print(wq.columns)
```

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
      'potential of hydrogen', 'sulphates', 'alcohol', 'quality', 'Id',
      'total_free_ratio'],
      dtype='object')
```

#12

```
aggrgatedata=wq.groupby('quality').mean()
aggrgatedata.head(5)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density
quality								
3	8.450000	0.897500	0.211667	2.666667	0.105333	8.166667	24.500000	0.997682
4	7.809091	0.700000	0.165758	2.566667	0.094788	14.848485	40.606061	0.996669
5	8.161077	0.585280	0.240124	2.540476	0.091770	16.612836	55.299172	0.997073
6	8.317749	0.504957	0.263680	2.444805	0.085281	15.215368	39.941558	0.996610

#13

```
slicedframe = (wq['total_free_ratio'] > 2.7) & (wq['total_free_ratio'] < 3.2)
```

```
sf=wq[slicedframe]
print(sf.head(5))
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
4	7.4	0.70	0.00	1.9	0.076	
5	7.4	0.66	0.00	1.8	0.075	
12	8.5	0.28	0.56	1.8	0.092	
14	7.6	0.39	0.31	2.3	0.082	

	free sulfur dioxide	total sulfur dioxide	density	potential of hydrogen	\
0	11.0	34.0	0.9978	3.51	
4	11.0	34.0	0.9978	3.51	
5	13.0	40.0	0.9978	3.51	
12	35.0	103.0	0.9969	3.30	
14	23.0	71.0	0.9982	3.52	

	sulphates	alcohol	quality	Id	total_free_ratio
0	0.56	9.4	5	0	3.090909
4	0.56	9.4	5	4	3.090909
5	0.56	9.4	5	5	3.076923
12	0.75	10.5	7	16	2.942857
14	0.65	9.7	5	21	3.086957

```
#14
wq.drop(wq.columns[2],axis=1)
```

	fixed acidity	volatile acidity	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	potential of hydrogen	sulph
0	7.4	0.700	1.9	0.076	11.0	34.0	0.99780	3.51	
1	7.8	0.880	2.6	0.098	25.0	67.0	0.99680	3.20	
2	7.8	0.760	2.3	0.092	15.0	54.0	0.99700	3.26	
3	11.2	0.280	1.9	0.075	17.0	60.0	0.99800	3.16	
4	7.4	0.700	1.9	0.076	11.0	34.0	0.99780	3.51	
...	...	...	...	...	...	...	...	...	
1138	6.3	0.510	2.3	0.076	29.0	40.0	0.99574	3.42	
1139	6.8	0.620	1.9	0.068	28.0	38.0	0.99651	3.42	
1140	6.2	0.600	2.0	0.090	32.0	44.0	0.99490	3.45	
1141	5.9	0.550	2.2	0.062	39.0	51.0	0.99512	3.52	
1142	5.9	0.645	2.0	0.075	32.0	44.0	0.99547	3.57	

```
#15
print(wq.columns[0])
print((wq.sort_values(by=wq.columns[0])).head(5))
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
32	4.6	0.52	0.15	2.1	0.054	
589	4.9	0.42	0.00	2.1	0.048	
396	5.0	1.04	0.24	1.6	0.050	
935	5.0	0.74	0.00	1.2	0.041	
898	5.0	0.38	0.01	1.6	0.048	

	free sulfur dioxide	total sulfur dioxide	density	\
32	8.0	65.0	0.99340	
589	16.0	42.0	0.99154	
396	32.0	96.0	0.99340	
935	16.0	46.0	0.99258	
898	26.0	60.0	0.99084	

	potential of hydrogen	sulphates	alcohol	quality	Id	\
32	3.90	0.56	13.1	4	45	
589	3.71	0.74	14.0	7	821	
396	3.74	0.62	11.5	5	553	
935	4.01	0.59	12.5	6	1321	
898	3.70	0.75	14.0	6	1270	

	total_free_ratio
32	8.125000



589	2.625000
396	3.000000
935	2.875000
898	2.307692

```
#15
print(wq.columns[0])
print((wq.sort_values(by=wq.columns[0],ascending=False)).head(5))
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
462	15.9	0.360	0.65	7.5	0.096	
311	15.6	0.685	0.76	3.7	0.100	
399	15.6	0.645	0.49	4.2	0.095	
397	15.5	0.645	0.49	4.2	0.095	
171	15.0	0.210	0.44	2.2	0.075	

	free sulfur dioxide	total sulfur dioxide	density	\
462	22.0	71.0	0.99760	
311	6.0	43.0	1.00320	
399	10.0	23.0	1.00315	
397	10.0	23.0	1.00315	
171	10.0	24.0	1.00005	

	potential of hydrogen	sulphates	alcohol	quality	Id	total_free_ratio
462	2.98	0.84	14.9	5	652	3.227273
311	2.95	0.68	11.2	7	442	7.166667
399	2.92	0.74	11.1	5	557	2.300000
397	2.92	0.74	11.1	5	554	2.300000
171	3.07	0.84	9.2	7	244	2.400000

```
#16
s=wq[wq.columns[2]].sum()
print(s)
```

306.74

```
#17
min=wq[wq.columns[3]].min()
max=wq[wq.columns[3]].max()
print(max,min)
```

15.5 0.9

```
#18
wq.rename(columns={"potential of hydrogen":"ph"},inplace=True)
print(wq.columns)
```

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
      'ph', 'sulphates', 'alcohol', 'quality', 'Id', 'total_free_ratio'],
      dtype='object')
```

```
#20
wq.to_csv("balance.csv",index=True,header=True);
```

```
#21
wq.fillna(0);
print(wq.isnull().sum());
```


fixed acidity	0
volatile acidity	0
citric acid	0
residual sugar	0
chlorides	0
free sulfur dioxide	0
total sulfur dioxide	0
density	0
ph	0
sulphates	0
alcohol	0
quality	0
Id	0
total_free_ratio	0
dtype: int64	

```
wq1=wq.iloc[:600];
```

```
wq2=wq.iloc[600:];
print(len(wq1))
print(len(wq2))
```

600  
543

```
index = wq[wq.columns[2]].idxmax()
row = df.iloc[index];
print(row)
```

	Pregnancies	1.000
	Glucose	81.000
	BloodPressure	72.000
	SkinThickness	18.000
	Insulin	40.000
	BMI	26.600
	DiabetesPedigreeFunction	0.283
	Age	24.000
	Outcome	0.000
	Name: 103, dtype: float64	